

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense, Activation
```

Task 1

Reading the dataset and preprocessing

```
#Reading the dataset
df = pd.read_csv('/content/drug200.csv')
df
```

| | Age | Sex | BP | Cholesterol | Na_to_K | Drug |
|-----|-----|-----|--------|-------------|---------|-------|
| 0 | 23 | F | HIGH | HIGH | 25.355 | drugY |
| 1 | 47 | M | LOW | HIGH | 13.093 | drugC |
| 2 | 47 | M | LOW | HIGH | 10.114 | drugC |
| 3 | 28 | F | NORMAL | HIGH | 7.798 | drugX |
| 4 | 61 | F | LOW | HIGH | 18.043 | drugY |
| ... | ... | ... | ... | ... | ... | ... |
| 195 | 56 | F | LOW | HIGH | 11.567 | drugC |
| 196 | 16 | M | LOW | HIGH | 12.006 | drugC |
| 197 | 52 | M | NORMAL | HIGH | 9.894 | drugX |
| 198 | 23 | M | NORMAL | NORMAL | 14.020 | drugX |
| 199 | 40 | F | LOW | NORMAL | 11.349 | drugX |

200 rows × 6 columns

```
#Data Preprocessing
df.isnull().any()
```

```
Age      False
Sex      False
BP       False
Cholesterol  False
Na_to_K  False
Drug     False
dtype: bool
```

```
df.isnull().sum()
```

```
Age      0
Sex      0
BP       0
Cholesterol  0
Na_to_K  0
Drug     0
dtype: int64
```

```
print(df.head())
```

| | Age | Sex | BP | Cholesterol | Na_to_K | Drug |
|---|-----|-----|--------|-------------|---------|-------|
| 0 | 23 | F | HIGH | HIGH | 25.355 | drugY |
| 1 | 47 | M | LOW | HIGH | 13.093 | drugC |
| 2 | 47 | M | LOW | HIGH | 10.114 | drugC |
| 3 | 28 | F | NORMAL | HIGH | 7.798 | drugX |
| 4 | 61 | F | LOW | HIGH | 18.043 | drugY |


```
df.describe(include = 'all')
```

| | Age | Sex | BP | Cholesterol | Na_to_K | Drug |
|--------|------------|-----|------|-------------|------------|-------|
| count | 200.000000 | 200 | 200 | 200 | 200.000000 | 200 |
| unique | NaN | 2 | 3 | 2 | NaN | 5 |
| top | NaN | M | HIGH | HIGH | NaN | drugY |
| freq | NaN | 104 | 77 | 103 | NaN | 91 |
| mean | 44.315000 | NaN | NaN | NaN | 16.084485 | NaN |
| std | 16.544315 | NaN | NaN | NaN | 7.223956 | NaN |
| min | 15.000000 | NaN | NaN | NaN | 6.269000 | NaN |
| 25% | 31.000000 | NaN | NaN | NaN | 10.445500 | NaN |
| 50% | 45.000000 | NaN | NaN | NaN | 13.936500 | NaN |
| 75% | 58.000000 | NaN | NaN | NaN | 16.280000 | NaN |

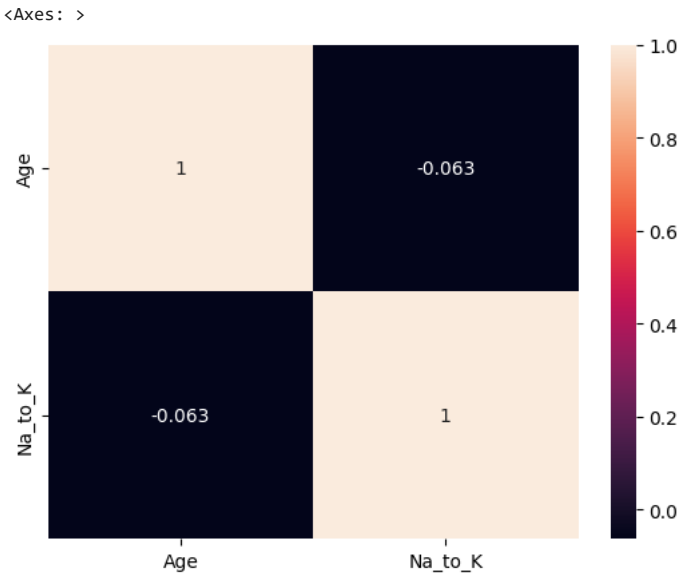
```
df['Drug'].value_counts()
```

```
drugY    91
drugX    54
drugA    23
drugC    16
drugB    16
Name: Drug, dtype: int64
```

```
#box plots
sns.boxplot(df)
```

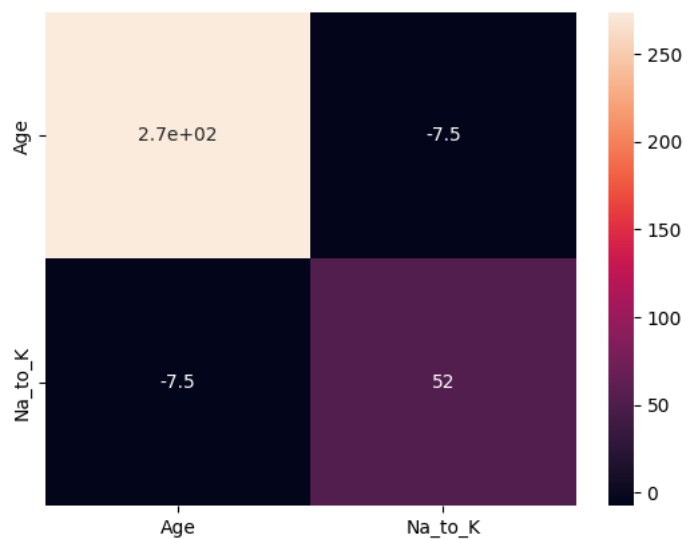


```
sns.heatmap(df.corr(),annot = True)
```




```
sns.heatmap(df.cov(),annot=True)
```

<Axes: >




```
sns.distplot(df['Na_to_K'])
```

<Axes: xlabel='Na_to_K', ylabel='Density'>



```
sns.lineplot(x = df['Age'],y=df['Na_to_K'])
```

<Axes: xlabel='Age', ylabel='Na_to_K'>



```

from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
df['Sex'] = label_encoder.fit_transform(df['Sex'])
df['BP'] = label_encoder.fit_transform(df['BP'])
df['Cholesterol'] = label_encoder.fit_transform(df['Cholesterol'])

# Normalize numerical variables
scaler = preprocessing.StandardScaler()
df[['Age', 'Na_to_K']] = scaler.fit_transform(df[['Age', 'Na_to_K']])

# Split the data into features and labels
x = df.iloc[:, :-1]
y = df.iloc[:, -1]
y = pd.get_dummies(y).values
print(df['Drug'].unique())

['drugY' 'drugC' 'drugX' 'drugA' 'drugB']

```

```

#split the data set into training and testing sets
xtrain,xtest,ytrain,ytest = train_test_split(x,y, test_size = 0.2, random_state=15)
print(x)
print(y)

```

| | Age | Sex | BP | Cholesterol | Na_to_K |
|-----|-----------|-----|----|-------------|-----------|
| 0 | -1.291591 | 0 | 0 | 0 | 1.286522 |
| 1 | 0.162699 | 1 | 1 | 0 | -0.415145 |
| 2 | 0.162699 | 1 | 1 | 0 | -0.828558 |
| 3 | -0.988614 | 0 | 2 | 0 | -1.149963 |
| 4 | 1.011034 | 0 | 1 | 0 | 0.271794 |
| .. | ... | ... | .. | ... | ... |
| 195 | 0.708057 | 0 | 1 | 0 | -0.626917 |
| 196 | -1.715759 | 1 | 1 | 0 | -0.565995 |
| 197 | 0.465676 | 1 | 2 | 0 | -0.859089 |
| 198 | -1.291591 | 1 | 2 | 1 | -0.286500 |
| 199 | -0.261469 | 0 | 1 | 1 | -0.657170 |

[200 rows x 5 columns]

```

[[0 0 0 0 1]
 [0 0 1 0 0]
 [0 0 1 0 0]
 [0 0 0 1 0]
 [0 0 0 0 1]
 [0 0 0 1 0]
 [0 0 0 0 1]
 [0 0 1 0 0]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 1 0 0]
 [0 0 0 0 1]
 [0 0 1 0 0]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 1 0]
 [0 0 0 0 1]
 [0 0 0 1 0]
 [0 0 0 0 0]
 [0 0 1 0 0]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 1 0]
 [0 0 0 0 1]
 [0 0 0 0 1]
 [0 0 0 1 0]
 [0 1 0 0 0]
 [0 0 0 1 0]
 [0 0 0 0 1]
 [0 0 0 1 0]
 [0 0 0 1 0]
 [0 0 0 1 0]
 [1 0 0 0 0]
 [0 0 0 1 0]
 [0 0 0 1 0]
 [0 0 0 1 0]
 [0 0 0 0 1]
 [0 1 0 0 0]
 [0 0 0 0 1]
 [0 0 0 1 0]

```

```
x.shape,y.shape
```

```
((200, 5), (200, 5))
```

xtrain

| | Age | Sex | BP | Cholesterol | Na_to_K |
|-----|-----------|-----|-----|-------------|-----------|
| 47 | 1.435202 | 1 | 1 | 0 | -0.803995 |
| 142 | 0.950439 | 1 | 0 | 1 | -1.035750 |
| 149 | -1.352186 | 1 | 1 | 0 | -1.100975 |
| 152 | 0.647462 | 1 | 2 | 1 | -1.224485 |
| 182 | -1.473377 | 0 | 1 | 1 | -0.610403 |
| ... | ... | ... | ... | ... | ... |
| 156 | -0.806828 | 1 | 0 | 1 | -0.674101 |
| 128 | 0.162699 | 1 | 1 | 1 | 2.422679 |
| 119 | 1.011034 | 0 | 0 | 0 | 1.303175 |
| 133 | -1.230996 | 1 | 2 | 0 | 1.346334 |
| 140 | 0.283889 | 1 | 0 | 1 | -1.362151 |

160 rows × 5 columns

Task 2

Build the ANN model

```
model = Sequential()
model.add(Dense(32, activation='relu', input_dim=x.shape[1]))#input layer
model.add(Dense(16, activation='relu')) #hidden layer
model.add(Dense(8, activation='relu')) #hidden layer
model.add(Dense(4, activation='relu')) #hidden layer
model.add(Dense(5, activation='softmax')) #Output layer
```

```
#Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense) | (None, 32) | 192 |
| dense_1 (Dense) | (None, 16) | 528 |
| dense_2 (Dense) | (None, 8) | 136 |
| dense_3 (Dense) | (None, 4) | 36 |
| dense_4 (Dense) | (None, 5) | 25 |

```
=====
Total params: 917
Trainable params: 917
Non-trainable params: 0
=====
```

```
#Training the model
model.fit(xtrain,ytrain,epochs=10,batch_size=5,validation_data=(xtest,ytest))
```

```
Epoch 1/10
32/32 [=====] - 1s 8ms/step - loss: 1.5833 - accuracy: 0.4500 - val_loss: 1.5270 - val_accuracy: 0.4750
Epoch 2/10
32/32 [=====] - 0s 2ms/step - loss: 1.4501 - accuracy: 0.4500 - val_loss: 1.4008 - val_accuracy: 0.4750
Epoch 3/10
32/32 [=====] - 0s 2ms/step - loss: 1.3154 - accuracy: 0.4500 - val_loss: 1.2479 - val_accuracy: 0.4750
Epoch 4/10
32/32 [=====] - 0s 2ms/step - loss: 1.1855 - accuracy: 0.4500 - val_loss: 1.1322 - val_accuracy: 0.4750
Epoch 5/10
32/32 [=====] - 0s 2ms/step - loss: 1.0947 - accuracy: 0.4812 - val_loss: 1.0402 - val_accuracy: 0.5000
Epoch 6/10
32/32 [=====] - 0s 2ms/step - loss: 1.0170 - accuracy: 0.5750 - val_loss: 0.9702 - val_accuracy: 0.5000
Epoch 7/10
32/32 [=====] - 0s 3ms/step - loss: 0.9555 - accuracy: 0.6000 - val_loss: 0.9188 - val_accuracy: 0.5250
Epoch 8/10
32/32 [=====] - 0s 3ms/step - loss: 0.9111 - accuracy: 0.6000 - val_loss: 0.8762 - val_accuracy: 0.5250
```

```
Epoch 9/10
32/32 [=====] - 0s 2ms/step - loss: 0.8675 - accuracy: 0.6062 - val_loss: 0.8407 - val_accuracy: 0.5250
Epoch 10/10
32/32 [=====] - 0s 3ms/step - loss: 0.8336 - accuracy: 0.5875 - val_loss: 0.8031 - val_accuracy: 0.5250
<keras.callbacks.History at 0x7f8441af760>
```

```
ypred = model.predict(xtest)
```

```
2/2 [=====] - 0s 8ms/step
```

```
ypred.shape
```

```
(40, 5)
```

▼ Task 3

Testing the data with the random values

```
ypred1 = model.predict([[1.432,0,0,1,-1.094]])
ypred1
```

```
1/1 [=====] - 0s 79ms/step
array([[0.19395398, 0.15538718, 0.20652266, 0.25762147, 0.18651469]],
      dtype=float32)
```